

# An Extensible Scheduler for the OpenLambda FaaS Platform

**Gustavo Totoy**, Edwin F. Boza, Cristina L. Abad



# Bio

- Game Developer at <http://freakycreations.net/>, a studio based in Ecuador




# Bio

- Game Developer at <http://freakycreations.net/>, a studio based in Ecuador
- Working on To Leave for its PS4 release




# Bio

- Game Developer at <http://freakycreations.net/>, a studio based in Ecuador
- Working on To Leave for its PS4 release
- Undergraduate student graduating in September 2018



# Bio

- Game Developer at <http://freakycreations.net/>, a studio based in Ecuador
- Working on To Leave for its PS4 release
- Undergraduate student graduating in September 2018
  - Applying for US grad school this year!





# Bio

- Game Developer at <http://freakycreations.net/>, a studio based in Ecuador
- Working on To Leave for its PS4 release
- Undergraduate student graduating in September 2018
  - Applying for US grad school this year!
- Research assistant at the Systems Research Lab lead by Prof. Cristina Abad in ESPOL University



# Bio

- Game Developer at <http://freakycreations.net/>, a studio based in Ecuador
- Working on To Leave for its PS4 release
- Undergraduate student graduating in September 2018
  - Applying for US grad school this year!
- Research assistant at the Systems Research Lab lead by Prof. Cristina Abad in ESPOL University
- Try to continue improving on my hobbies: Cook, CrossFit and BJJ

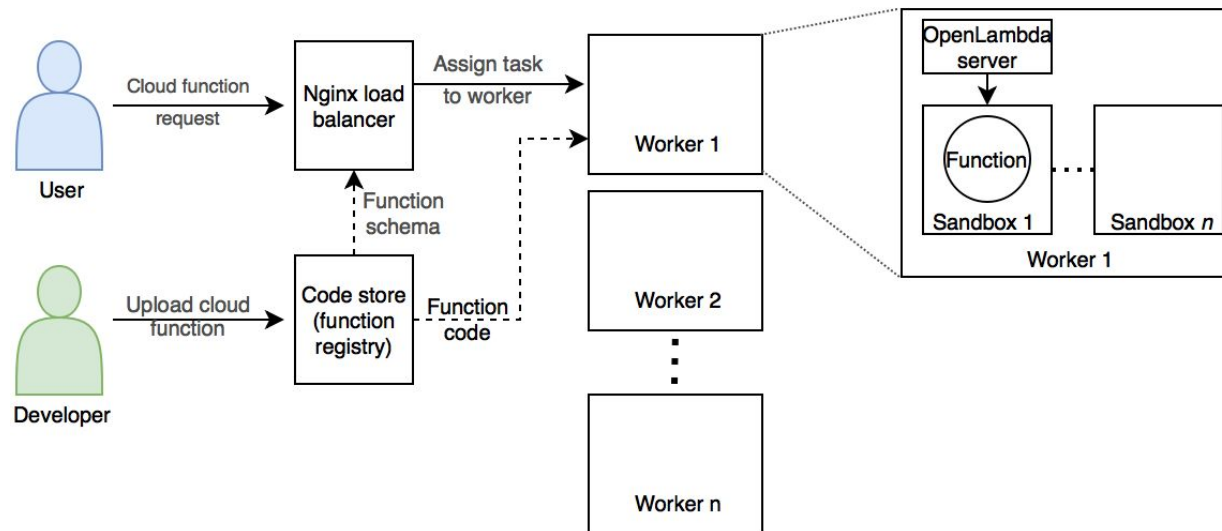


olscheduler



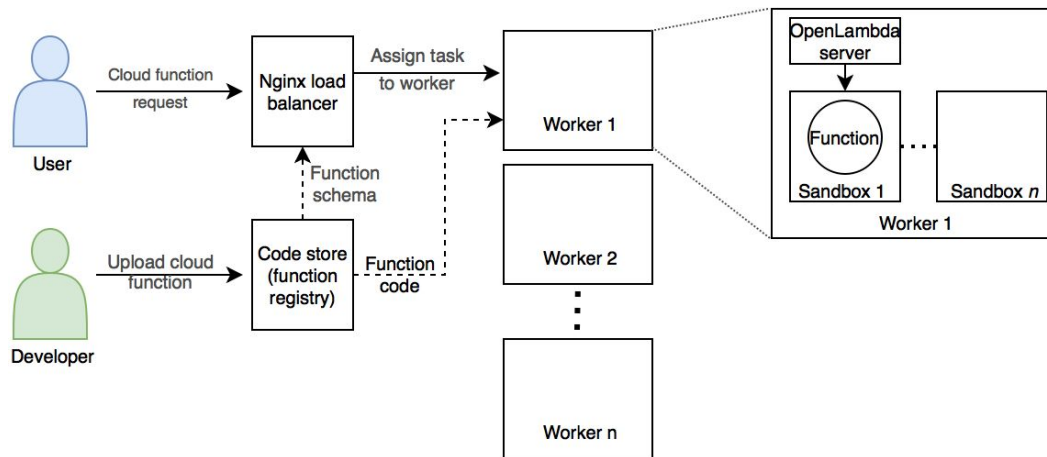
# Summary

- Simple and extensible function scheduler for the OpenLambda Function as a Service platform



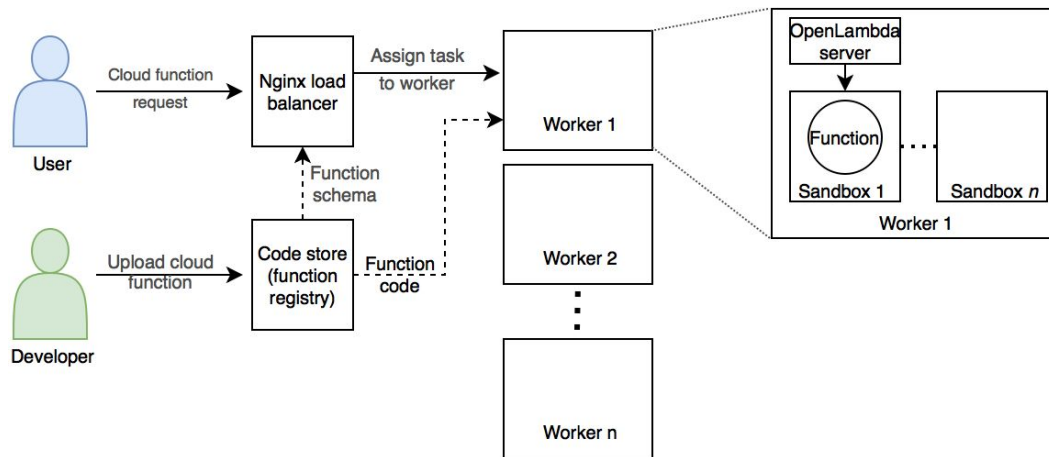
# Summary

- Simple and extensible function scheduler for the OpenLambda Function as a Service platform
- Written in Go



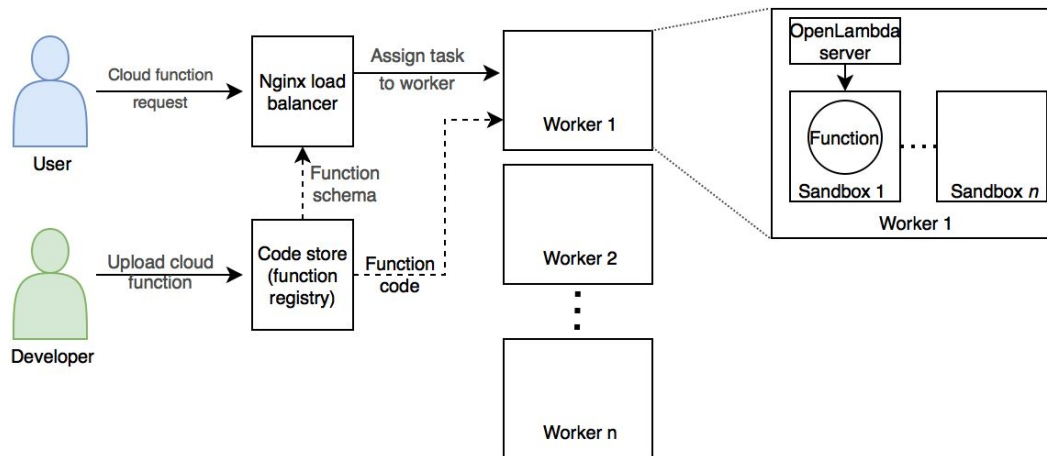
# Summary

- Simple and extensible function scheduler for the OpenLambda Function as a Service platform
- Written in Go
- ~400 lines of code



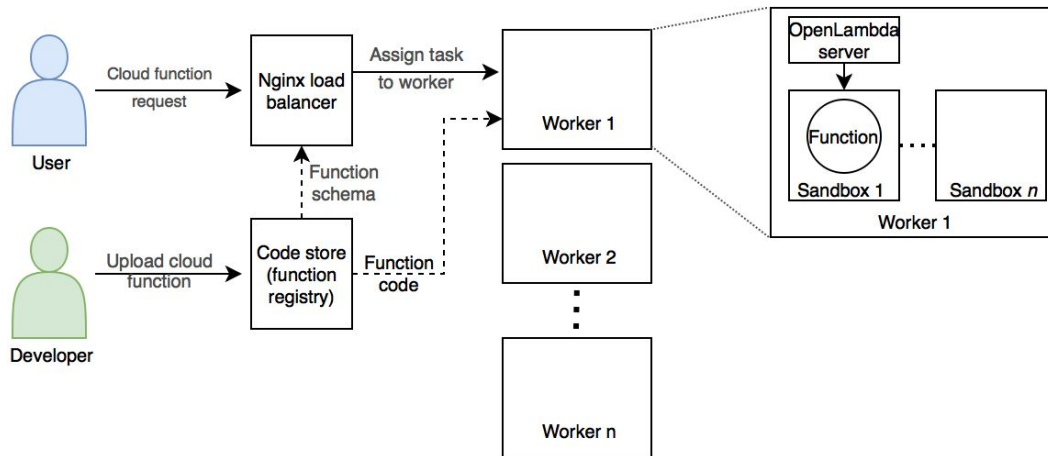
# Summary

- Simple and extensible function scheduler for the OpenLambda Function as a Service platform
- Written in Go
- ~400 lines of code
- Uses Go standard library



# Summary

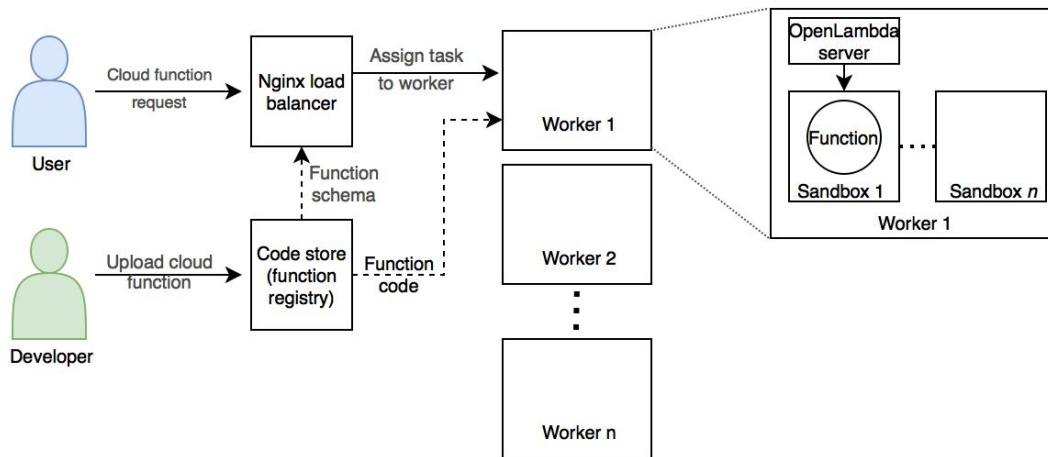
- Simple and extensible function scheduler for the OpenLambda Function as a Service platform
- Written in Go
- ~400 lines of code
- Uses Go standard library
- Offers four scheduling policies: random, round-robin, least-loaded and pkg-aware





# Summary

- Simple and extensible function scheduler for the OpenLambda Function as a Service platform
- Written in Go
- ~400 lines of code
- Uses Go standard library
- Offers four scheduling policies: random, round-robin, least-loaded and pkg-aware
- Open Source:  
<https://github.com/gtotoy/olscheduler>



# Motivation

Load Balancer -> OpenLambda + PipSqueak

# Background

- A Function-as-a-Service (FaaS) cloud platforms enable tenants to deploy and execute functions on the cloud, without having to worry about server provisioning

# Background

- A Function-as-a-Service (FaaS) cloud platforms enable tenants to deploy and execute functions on the cloud, without having to worry about server provisioning
- Small cloud functions with big or many packages to be loaded

# Background

- A Function-as-a-Service (FaaS) cloud platforms enable tenants to deploy and execute functions on the cloud, without having to worry about server provisioning
- Small cloud functions with big or many packages to be loaded
- In a distributed computing platform, co-locating tasks at worker nodes that cache any required files is a time-proven mechanism to reduce task latency

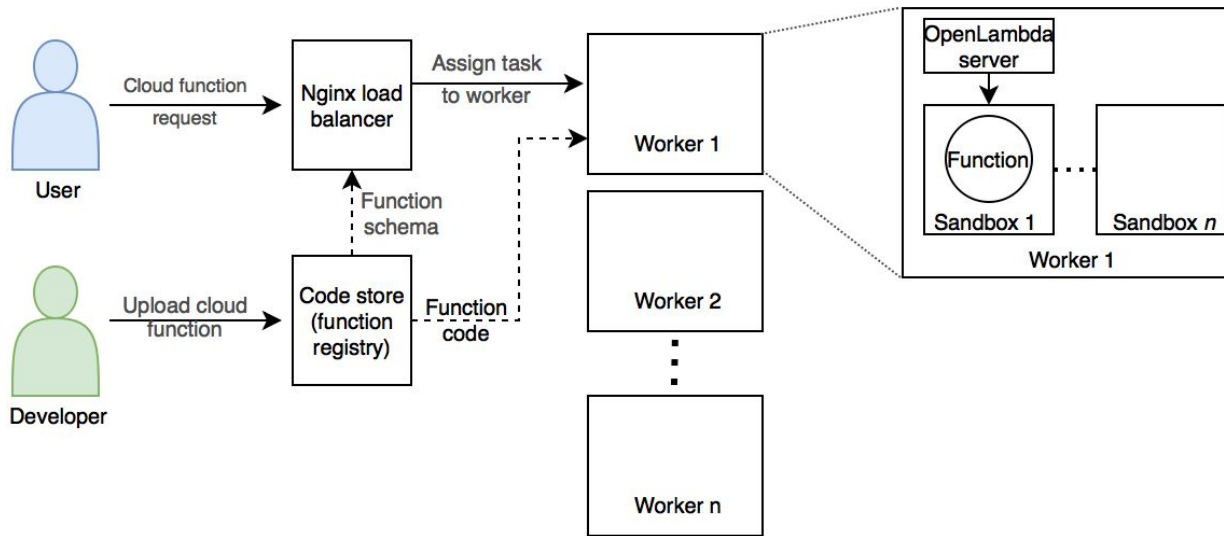


# Background

- A Function-as-a-Service (FaaS) cloud platforms enable tenants to deploy and execute functions on the cloud, without having to worry about server provisioning
- Small cloud functions with big or many packages to be loaded
- In a distributed computing platform, co-locating tasks at worker nodes that cache any required files is a time-proven mechanism to reduce task latency
- We seek to increase the hit rate of the package cache and, as a result, reduce the latency of the cloud functions

# Load Balancer -> *OpenLambda* + PipSqueak


- OpenLambda [1]: is a serverless computing platform that supports the Function as a Service execution model



[1] Hendrickson, S., Sturdevant, S., Harter, T., Venkataramani, V., ArpaciDusseau, A., and Arpaci-Dusseau, R. Serverless computation with OpenLambda. In USENIX Work. Hot Topics in Cloud Comp. (HotCloud) (2016).

# Load Balancer -> OpenLambda + *PipSqueak*

- Pipsqueak [2]: shared package cache available at each OpenLambda worker
- Stores a set of Python interpreters with pre-imported packages, in a sleeping state



# ***Load Balancer*** -> OpenLambda + PipSqueak

- Nginx ([www.nginx.com](http://www.nginx.com)), which comes with load balancing methods: round robin, least connected and ip hash
- gobetween ([gobetween.io](http://gobetween.io))
- But, What if we want to make smarter scheduling policies?

# olscheduler

Our scheduler exposes functions and data structures that can be used to get the following information:

1. **Workers:** Number and references to the worker lambda nodes
2. **Per-worker load:** Measured as the number of active requests that each worker is currently handling
3. **Required packages:** The list of required packages, sorted by size, exposed for each function call



# olscheduler

Currently supports the following scheduling algorithms:

- **Round-robin**

# olscheduler

Currently supports the following scheduling algorithms:

- **Round-robin**
- **Least-loaded**

# olscheduler

Currently supports the following scheduling algorithms:

- **Round-robin**
- **Least-loaded**
- **Random**

# olscheduler

Currently supports the following scheduling algorithms:

- **Round-robin**
- **Least-loaded**
- **Random**
- **Pkg-aware:** Seeks to maximize cache affinity (with respect to the packages in the package cache), while avoiding overloading workers beyond a configurable threshold

# olscheduler

- Small code base ~400 LOCs

# olscheduler

- Small code base ~400 LOCs
- Simpler to modify and extend than the nginx load balancer

# olscheduler

- Small code base ~400 LOCs
- Simpler to modify and extend than the nginx load balancer
- Written in a modern programming language

# olscheduler

- Small code base ~400 LOCs
- Simpler to modify and extend than the nginx load balancer
- Written in a modern programming language
- It can be adapted to OpenLambda specific requirements



# olscheduler

- Small code base ~400 LOCs
- Simpler to modify and extend than the nginx load balancer
- Written in a modern programming language
- It can be adapted to OpenLambda specific requirements
- Offers pkg-aware, a novel scheduler policy

# olscheduler

- Small code base ~400 LOCs
- Simpler to modify and extend than the nginx load balancer
- Written in a modern programming language
- It can be adapted to OpenLambda specific requirements
- Offers pkg-aware, a novel scheduler policy
- Allows development of additional scheduling policies

# olscheduler

- Small code base ~400 LOCs
- Simpler to modify and extend than the nginx load balancer
- Written in a modern programming language
- It can be adapted to OpenLambda specific requirements
- Offers pkg-aware, a novel scheduler policy
- Allows development of additional scheduling policies
- Facilitates experiment reproducibility

# Pkg-Aware Algorithm

Abad, C. L., Boza, E. F., and van Eyk, E. Package-aware scheduling of FaaS functions. In HotCloudPerf workshop, co-located with ACM/SPEC Intl. Conf. Perf. Eng. (ICPE) (2018)

# Algorithm

- The algorithm seeks to maximize cache affinity (with respect to the packages in the package cache), while avoiding overloading workers beyond a configurable threshold

**Algorithm 1:** Package-aware scheduler algorithm for Open-Lambda

**Global data:** List of workers,  $W = w_1, \dots, w_n$ , Hash functions  $H_1$  and  $H_2$ , maximum load threshold,  $t$

**Input:** Function,  $f$ , list of required packages sorted by descending package size,  $P = p_1, \dots, p_n$

```
1 if ( $P$  is not empty)then
    /* Greedily seek affinity w/ large package */
2   for ( $l = 1, \dots, |P|$ )do
    /* Calculate two possible worker targets */
3      $t1 = H_1(p_l) \% |W| + 1$ 
4      $t2 = H_2(p_l) \% |W| + 1$ 
    /* Select target with least load */
5     if ( $load(w_{t1}) < load(w_{t2})$ )then
6        $A := t1$ 
7     else
8        $A := t2$ 
    /* If target is not overloaded, we are done */
9     if ( $load(w_A) < t$ )then
10      Assign  $f$  to  $w_A$ 
11      return
    /* Balance load */
12 Assign  $f$  to least loaded worker,  $w_i$ 
```

# Algorithm

- The algorithm seeks to maximize cache affinity (with respect to the packages in the package cache), while avoiding overloading workers beyond a configurable threshold

**Algorithm 1:** Package-aware scheduler algorithm for Open-Lambda

**Global data:** List of workers,  $W = w_1, \dots, w_n$ , Hash functions  $H_1$  and  $H_2$ , maximum load threshold,  $t$

**Input:** Function,  $f$ , list of required packages sorted by descending package size,  $P = p_1, \dots, p_n$

```
1 if ( $P$  is not empty)then
    /* Greedily seek affinity w/ large package */
2   for ( $l = 1, \dots, |P|$ )do
    /* Calculate two possible worker targets */
3      $t1 = H_1(p_l) \% |W| + 1$ 
4      $t2 = H_2(p_l) \% |W| + 1$ 
    /* Select target with least load */
5     if ( $load(w_{t1}) < load(w_{t2})$ )then
6        $A := t1$ 
7     else
8        $A := t2$ 
    /* If target is not overloaded, we are done */
9     if ( $load(w_A) < t$ )then
10      Assign  $f$  to  $w_A$ 
11      return
    /* Balance load */
12 Assign  $f$  to least loaded worker,  $w_i$ 
```

# Algorithm

- The algorithm seeks to maximize cache affinity (with respect to the packages in the package cache), while avoiding overloading workers beyond a configurable threshold

**Algorithm 1:** Package-aware scheduler algorithm for Open-Lambda

**Global data:** List of workers,  $W = w_1, \dots, w_n$ , Hash functions  $H_1$  and  $H_2$ , maximum load threshold,  $t$

**Input:** Function,  $f$ , list of required packages sorted by descending package size,  $P = p_1, \dots, p_n$

```
1 if ( $P$  is not empty)then
    /* Greedily seek affinity w/ large package */
2   for ( $l = 1, \dots, |P|$ )do
    /* Calculate two possible worker targets */
3      $t1 = H_1(p_l) \% |W| + 1$ 
4      $t2 = H_2(p_l) \% |W| + 1$ 
    /* Select target with least load */
5     if ( $load(w_{t1}) < load(w_{t2})$ )then
6        $A := t1$ 
7     else
8        $A := t2$ 
    /* If target is not overloaded, we are done */
9     if ( $load(w_A) < t$ )then
10      Assign  $f$  to  $w_A$ 
11      return
    /* Balance load */
12 Assign  $f$  to least loaded worker,  $w_i$ 
```

# Algorithm

- The algorithm seeks to maximize cache affinity (with respect to the packages in the package cache), while avoiding overloading workers beyond a configurable threshold

**Algorithm 1:** Package-aware scheduler algorithm for Open-Lambda

**Global data:** List of workers,  $W = w_1, \dots, w_n$ , Hash functions  $H_1$  and  $H_2$ , maximum load threshold,  $t$

**Input:** Function,  $f$ , list of required packages sorted by descending package size,  $P = p_1, \dots, p_n$

```
1 if ( $P$  is not empty)then
    /* Greedily seek affinity w/ large package */
2   for ( $l = 1, \dots, |P|$ )do
    /* Calculate two possible worker targets */
3      $t1 = H_1(p_l) \% |W| + 1$ 
4      $t2 = H_2(p_l) \% |W| + 1$ 
    /* Select target with least load */
5     if ( $load(w_{t1}) < load(w_{t2})$ )then
6        $A := t1$ 
7     else
8        $A := t2$ 
    /* If target is not overloaded, we are done */
9     if ( $load(w_A) < t$ )then
10      Assign  $f$  to  $w_A$ 
11      return
    /* Balance load */
12 Assign  $f$  to least loaded worker,  $w_i$ 
```



# Algorithm

- The algorithm seeks to maximize cache affinity (with respect to the packages in the package cache), while avoiding overloading workers beyond a configurable threshold

**Algorithm 1:** Package-aware scheduler algorithm for Open-Lambda

**Global data:** List of workers,  $W = w_1, \dots, w_n$ , Hash functions  $H_1$  and  $H_2$ , maximum load threshold,  $t$

**Input:** Function,  $f$ , list of required packages sorted by descending package size,  $P = p_1, \dots, p_n$

```
1 if ( $P$  is not empty)then
    /* Greedily seek affinity w/ large package */
2   for ( $l = 1, \dots, |P|$ )do
    /* Calculate two possible worker targets */
3      $t1 = H_1(p_l) \% |W| + 1$ 
4      $t2 = H_2(p_l) \% |W| + 1$ 
    /* Select target with least load */
5     if ( $load(w_{t1}) < load(w_{t2})$ )then
6        $A := t1$ 
7     else
8        $A := t2$ 
    /* If target is not overloaded, we are done */
9     if ( $load(w_A) < t$ )then
10      Assign  $f$  to  $w_A$ 
11      return
    /* Balance load */
12  Assign  $f$  to least loaded worker,  $w_i$ 
```

\*/

# Algorithm

- The algorithm seeks to maximize cache affinity (with respect to the packages in the package cache), while avoiding overloading workers beyond a configurable threshold

**Algorithm 1:** Package-aware scheduler algorithm for Open-Lambda

**Global data:** List of workers,  $W = w_1, \dots, w_n$ , Hash functions  $H_1$  and  $H_2$ , maximum load threshold,  $t$

**Input:** Function,  $f$ , list of required packages sorted by descending package size,  $P = p_1, \dots, p_n$

```
1 if ( $P$  is not empty)then
    /* Greedily seek affinity w/ large package */
2   for ( $l = 1, \dots, |P|$ )do
    /* Calculate two possible worker targets */
3      $t1 = H_1(p_l) \% |W| + 1$ 
4      $t2 = H_2(p_l) \% |W| + 1$ 
    /* Select target with least load */
5     if ( $load(w_{t1}) < load(w_{t2})$ )then
6        $A := t1$ 
7     else
8        $A := t2$ 
    /* If target is not overloaded, we are done */
9     if ( $load(w_A) < t$ )then
10      Assign  $f$  to  $w_A$ 
11      return
    /* Balance load */
12 Assign  $f$  to least loaded worker,  $w_i$ 
```

# Algorithm

- The algorithm seeks to maximize cache affinity (with respect to the packages in the package cache), while avoiding overloading workers beyond a configurable threshold
- With the information exposed by olscheduler, implementing this policy was straightforward, and took only 46 LOCs

**Algorithm 1:** Package-aware scheduler algorithm for Open-Lambda

**Global data:** List of workers,  $W = w_1, \dots, w_n$ , Hash functions  $H_1$  and  $H_2$ , maximum load threshold,  $t$

**Input:** Function,  $f$ , list of required packages sorted by descending package size,  $P = p_1, \dots, p_n$

```
1 if ( $P$  is not empty)then
    /* Greedily seek affinity w/ large package */
2   for ( $l = 1, \dots, |P|$ )do
    /* Calculate two possible worker targets */
3      $t1 = H_1(p_l) \% |W| + 1$ 
4      $t2 = H_2(p_l) \% |W| + 1$ 
    /* Select target with least load */
5     if ( $load(w_{t1}) < load(w_{t2})$ )then
6        $A := t1$ 
7     else
8        $A := t2$ 
    /* If target is not overloaded, we are done */
9     if ( $load(w_A) < t$ )then
10      Assign  $f$  to  $w_A$ 
11      return
    /* Balance load */
12 Assign  $f$  to least loaded worker,  $w_i$ 
```

# 65%

Preliminary simulation results show that this simple approach can cut the function latency by more than 65%

# Results

- Improved median hit rate from 51.15% (least-loaded) to 63.52% (pkg-aware)

# Results

- Improved median hit rate from 51.15% (least-loaded) to 63.52% (pkg-aware)
- Median latency improves by 65.8% (pkg-aware vs. least-loaded), and tail latency improves by 41.9% (90th percentile)

# Results

- Improved median hit rate from 51.15% (least-loaded) to 63.52% (pkg-aware)
- Median latency improves by 65.8% (pkg-aware vs. least-loaded), and tail latency improves by 41.9% (90th percentile)
- If we compare against a least-loaded load balancer in an unoptimized platform that does not cache function packages, our algorithm improves median latency by 189.9 times

# Future Work

Real Cloud Experiments + More  
scheduling algorithms





# Planning

- PipSqueak has recently released a benchmarking suite for OpenLambda to evaluate its cache performance improvements
- We are going to use this suite to measure latency when pkg-aware scheduling policy is selected and compare against other available policies
- We plan to implement more refined version of pkg-aware based on the work of Package-aware scheduling of FaaS functions [3]

# Thanks!

Contact:

Systems Research Lab  
ESPOL University  
Guayaquil, Ecuador

[gtotoy@fiec.espol.edu.ec](mailto:gtotoy@fiec.espol.edu.ec)  
[cabad@fiec.espol.edu.ec](mailto:cabad@fiec.espol.edu.ec)

